

Damage Detection of Cars and Repair cost Estimation!!

It's more than just image segmentation



Using the best image segmentation model for instance segmentation to detect damaged areas on cars is not new and has been done before and here I'm also doing the same thing but with an extension of estimating repair costs.

Therefore I intend to visit the instance segmentation part in brief and discuss the repair costs estimation and future work in detail.

The Business Problem

Claiming Insurance for your damaged vehicle is a hassle not only for the insured but for the insurer too. It involves paperwork and inspection that consumes time, yet it is necessary to avoid fraudulent claims. The whole process takes at least a week to finalize the repair cost, after inspection, from the mechanic.

What if we estimate the repair cost of damages from pictures of the damaged areas through machine learning using the past claims data for training that the insurance company already has. This will reduce inspection time and consequently claims can be processed faster.

Though it is absolutely **necessary** to perform thorough inspection for **severe** damages where **internal** parts are affected, **minor/moderate** damages that has only affected the **external** surface of the vehicle can avoid it by estimating the repair cost and save a lot of time both for the insured and the insurer.

Dataset

I got the dataset from Kaggle Datasets [here](#). It contains total 78 images of damaged cars of which 59 are for training, 11 for validation and remaining 8 for testing.

There are two types of annotations per image:

1. All the damaged areas are annotated, with a label **damage**

2. All the parts of car are annotated with labels - **front bumper, rear bumper, headlamp, door** and **hood**.





Annotations of

damaged area (left) and car parts (right)

Modifying the damage annotations

Now, to help train my repair cost estimation model I modified the **damage annotations** into three categories: **minor**, **moderate** and **severe**.

It was a manual process where I looked at each car image with damages and edited the annotation json file into 3 categories using the following method:

- Scratches - classified as **minor** damage.
- Dents - classified as **moderate** damage.
- Broken part - classified as **severe** damage.

Now after this modification the updated annotations look like:



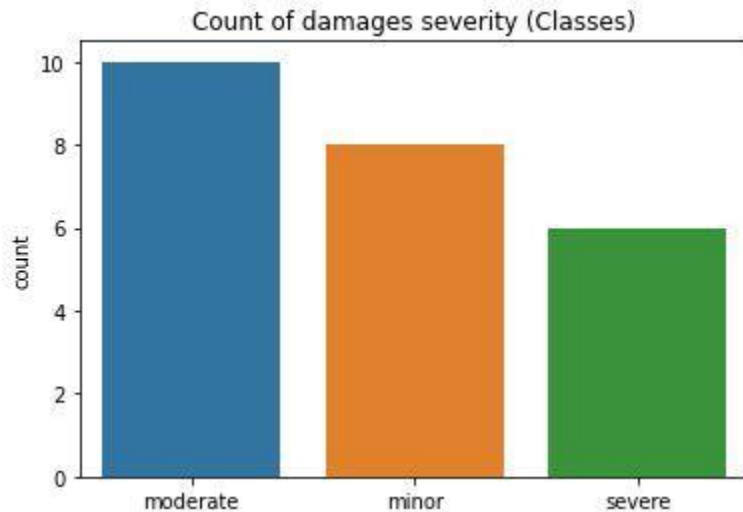
Annotations of damage type (image by author)

EDA

Just to recap, now I have two annotations in total, one for **damage types** (modified above) and another for **parts of car**.


Proper EDA gave me the following results:

- Damage type annotations have **moderate** damages the most followed by minor and severe damages.



Count of damage classes in train data

- Car parts annotations have the **headlamp** annotations the most and **rear bumper** annotations the least.



Count of cart part classes in train data

- All car parts are not annotated, E.g.- **front fender, rear fender, trunk, windshield** etc. Hence the image segmentation model won't be able to detect those parts.
- Training data has only 59 images, which is not enough, so we need a proper pretrained backbone architecture for image segmentation model (Mask R-CNN) to facilitate transfer learning.


Creating the training data

Just a brief introduction to the one of the most important feature in our training data, **dice coefficient**.

Dice Coefficient

Dice Coefficient is the area of overlap between two masks.

$DSC = 2 * (\text{the Area of Overlap}) / (\text{the total number of pixels in both masks})$



Formula for dice coefficient

Now using the above formula and the annotations we are provided with, we can find ***how much percentage of area is damaged in a particular car part***

Ex: In the above formula, let **A** be segmentation mask of hood of the car, and **B** be segmentation mask of moderate damaged area in the hood, then putting the values in the formula we get:

$DSC = 2 * (\text{the Area of Overlap between hood mask and moderate damage mask}) / (\text{the total number of pixels in both masks}) = 0.2345$
(let)

then we can say that **23%** of the hood is **moderately** damaged and '**hood_dice**' feature gets a value of 0.2345.

Similarly we can do this iteratively for every damage mask in the car image. The final training dataset looked something like this:

mp_dice	rear_bump_dice	door_dice	hood_dice	front_bump_dice	unknown	minor	moderate	severe	price	
1	0.0	0.000000	0.004600	0.000000	0.000000	0	0	1	0	500
2	0.0	0.000000	0.002410	0.000000	0.000000	0	0	1	0	50
3	0.0	0.041677	0.000000	0.000000	0.000000	0	1	0	0	400
4	0.0	0.026820	0.000000	0.000000	0.000000	0	1	0	0	400
5	0.0	0.000000	0.000000	0.000000	0.000000	1	0	1	0	250
6	0.0	0.000000	0.000000	0.001841	0.070522	0	0	0	1	1700

Modelling Approach

Intuitively there are total 3 models that we need to train:


1. First model to train the above dataset for estimating the price, lets call it **model 1**.
2. Second for detecting the **damaged areas** with their classes (**minor, moderate** and **severe**), lets call it **model 2**.
3. Third for detecting the **parts of car** with their classes (**front bumper, rear bumper, headlamp, door** and **hood**), lets call it **model 3**.

Model **1** can be trained using a basic logistic regression model, but unfortunately we have very less data for training and hence I would be calculating **mean** price using **group by** from identical rows in test data.

We need an Image segmentation model for **second** and **third** models to help generate the test data.

Mask R-CNN

There are many models available for instance segmentation. Here I have used **Mask R-CNN** which gives decent results considering the size of dataset we have.



Architecture of Mask R-CNN

The architecture contains the following components:

- **Backbone:** This is a standard convolutional neural network (typically, ResNet50 or ResNet101) that serves as a feature extractor.
- **Region Proposal Network (RPN):** It is a lightweight neural network that scans the image in a sliding-window fashion and finds areas that contain objects.
- **ROI Classifier:** This network is deeper and has the capacity to classify ROIs to specific classes (person, car, chair, etc.).
- **Bounding Box Regressor:** Very similar to how it's done in the RPN, and its purpose is to further refine the location and size of the bounding box to encapsulate the object.
- **ROI Pooling:** ROI pooling refers to cropping a part of a feature map and resizing it to a fixed size. The authors suggest a method they named **ROIAlign**, in which they

sample the feature map at different points and apply a bilinear interpolation.

- **Segmentation Masks:** The mask branch is a convolutional network that takes the positive regions selected by the ROI classifier and generates masks for them.

I used a pre-built library called **Detectron2** to train the Mask R-CNN Model.

Model Training and Performance

1. **Damage type detection:** Model **2** gave me an Average Precision (AP) of 8.89 over all classes for segmentation masks. See notebook snippet below:

Training output of Model 2 (damage type detection)





Predictions of damage type detection model

2. Car part detection: Model **3** gave me an Average Precision (AP) of 34.18 over all classes for segmentation masks. See notebook snippet below:

Training output of Model 3 (car part detection)

```
register_coco_instances("car_part_train", {},
os.path.join(COCO_DIR,train_dir,"COCO_mul_train_annos.json"),
os.path.join(COCO_DIR,train_dir))
register_coco_instances("car_part_val", {},
os.path.join(COCO_DIR,val_dir,"COCO_mul_val_annos.json"),
os.path.join(COCO_DIR,val_dir))

car_part_dataset_dicts = DatasetCatalog.get("car_part_train")
car_part_metadata_dicts = MetadataCatalog.get("car_part_train")
```

In []:


```

cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-
InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml"))
#cfg.merge_from_file('configs/COCO-
InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml')
cfg.DATASETS.TRAIN = ("car_part_train",)
cfg.DATASETS.TEST = ("car_part_val",)
cfg.INPUT.CROP.ENABLED = True
cfg.DATALOADER.NUM_WORKERS = 1
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-
InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml") # Let training initialize
from model zoo
cfg.SOLVER.IMS_PER_BATCH = 1
cfg.SOLVER.BASE_LR = 0.001 # pick a good LR
cfg.SOLVER.WARMUP_ITERS = 800
cfg.SOLVER.MAX_ITER = 2500 #adjust up if val mAP is still rising, adjust down
if overfit
cfg.SOLVER.STEPS = (600, 2000,)
cfg.SOLVER.GAMMA = 0.05
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 256 # faster, and good enough
for this dataset (default: 512)
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 5 # only has one class (damage)
cfg.MODEL.RETINANET.NUM_CLASSES = 5 # only has one class (damage)
cfg.TEST.EVAL_PERIOD = 400
cfg.SOLVER.CHECKPOINT_PERIOD = 400

```

```

# Clear any logs from previous runs
#TODO add timestamp to logs
!rm -rf cfg.OUTPUT_DIR

```

```

os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = CocoTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()

```

```
[05/19 16:17:44 d2.engine.defaults]: Model:
```

```

GeneralizedRCNN(
  (backbone): FPN(
    (fpn_lateral2): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (fpn_lateral3): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (fpn_lateral4): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))
    (fpn_lateral5): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), paddin
g=(1, 1))

```

```

(top_block): LastLevelMaxPool()
(bottom_up): ResNet(
  (stem): BasicStem(
    (conv1): Conv2d(
      3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=Fals
e
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
  )
  (res2): Sequential(
    (0): BottleneckBlock(
      (shortcut): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv1): Conv2d(
        64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=F
else
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
    )
    (1): BottleneckBlock(
      (conv1): Conv2d(
        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=F
else
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
    )
    (2): BottleneckBlock(
      (conv1): Conv2d(
        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=F
else
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)

```

```

    )
    (conv3): Conv2d(
      64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
  )
)
(res3): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv1): Conv2d(
      256, 128, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv2): Conv2d(
      128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv3): Conv2d(
      128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
  )
  (1): BottleneckBlock(
    (conv1): Conv2d(
      512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv2): Conv2d(
      128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv3): Conv2d(
      128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
  )
  (2): BottleneckBlock(
    (conv1): Conv2d(
      512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv2): Conv2d(
      128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv3): Conv2d(

```

```

        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
)
(3): BottleneckBlock(
  (conv1): Conv2d(
    512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv2): Conv2d(
    128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv3): Conv2d(
    128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
  )
)
)
(res4): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
    (conv1): Conv2d(
      512, 256, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (1): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)

```

```

    )
)
(2): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(3): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(4): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(5): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(

```

```

        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
)
(res5): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
    (conv1): Conv2d(
      1024, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
      512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
      512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
  )
  (1): BottleneckBlock(
    (conv1): Conv2d(
      2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
      512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
      512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
  )
  (2): BottleneckBlock(
    (conv1): Conv2d(
      2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
      512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False

```



```
(norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
      512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
  )
)
)
)
(proposal_generator): RPN(
  (rpn_head): StandardRPNHead(
    (conv): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
      (activation): ReLU()
    )
    (objectness_logits): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
    (anchor_deltas): Conv2d(256, 12, kernel_size=(1, 1), stride=(1, 1))
  )
  (anchor_generator): DefaultAnchorGenerator(
    (cell_anchors): BufferList()
  )
)
)
(roi_heads): StandardROIHeads(
  (box_pooler): ROIPooler(
    (level_poolers): ModuleList(
      (0): ROIAlign(output_size=(7, 7), spatial_scale=0.25, sampling_ratio=0, aligned=True)
      (1): ROIAlign(output_size=(7, 7), spatial_scale=0.125, sampling_ratio=0, aligned=True)
      (2): ROIAlign(output_size=(7, 7), spatial_scale=0.0625, sampling_ratio=0, aligned=True)
      (3): ROIAlign(output_size=(7, 7), spatial_scale=0.03125, sampling_ratio=0, aligned=True)
    )
  )
  (box_head): FastRCNNConvFCHead(
    (flatten): Flatten(start_dim=1, end_dim=-1)
    (fc1): Linear(in_features=12544, out_features=1024, bias=True)
    (fc_relu1): ReLU()
    (fc2): Linear(in_features=1024, out_features=1024, bias=True)
    (fc_relu2): ReLU()
  )
  (box_predictor): FastRCNNOutputLayers(
    (cls_score): Linear(in_features=1024, out_features=6, bias=True)
    (bbox_pred): Linear(in_features=1024, out_features=20, bias=True)
  )
  (mask_pooler): ROIPooler(
    (level_poolers): ModuleList(
      (0): ROIAlign(output_size=(14, 14), spatial_scale=0.25, sampling_ratio=0, aligned=True)
      (1): ROIAlign(output_size=(14, 14), spatial_scale=0.125, sampling_ratio=0, aligned=True)
    )
  )
)
```

```

        (2): ROIAlign(output_size=(14, 14), spatial_scale=0.0625, sampling_ratio=0, aligned=True)
        (3): ROIAlign(output_size=(14, 14), spatial_scale=0.03125, sampling_ratio=0, aligned=True)
    )
)
(mask_head): MaskRCNNConvUpsampleHead(
  (mask_fcn1): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
    (activation): ReLU()
  )
  (mask_fcn2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
    (activation): ReLU()
  )
  (mask_fcn3): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
    (activation): ReLU()
  )
  (mask_fcn4): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)
    (activation): ReLU()
  )
  (deconv): ConvTranspose2d(256, 256, kernel_size=(2, 2), stride=(2, 2))
  (deconv_relu): ReLU()
  (predictor): Conv2d(256, 5, kernel_size=(1, 1), stride=(1, 1))
)
)
)
[05/19 16:17:44 d2.data.datasets.coco]: Loaded 59 images in COCO format from
damage_dataset/train/COCO_mul_train_annos.json
[05/19 16:17:44 d2.data.build]: Removed 0 images with no usable annotations.
59 images left.
[05/19 16:17:44 d2.data.build]: Distribution of instances among all 5 categories:
| category | #instances | category | #instances | category | #instances |
|:-----|:-----|:-----|:-----|:-----|:-----|
| headlamp | 64 | rear_bumper | 24 | door | 32 |
| hood | 26 | front_bumper | 31 | | |
| total | 177 | | | | |
[05/19 16:17:44 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used in training: [RandomCrop(crop_type='relative_range', crop_size=[0.9, 0.9]), ResizeShortestEdge(short_edge_length=(640, 672, 704, 736, 768, 800), max_size=1333, sample_style='choice'), RandomFlip()]
[05/19 16:17:44 d2.data.build]: Using training sampler TrainingSampler
[05/19 16:17:44 d2.data.common]: Serializing 59 elements to byte tensors and concatenating them all ...
[05/19 16:17:44 d2.data.common]: Serialized dataset takes 0.03 MiB

```

Skip loading parameter 'roi_heads.box_predictor.cls_score.weight' to the model due to incompatible shapes: (81, 1024) in the checkpoint but (6, 1024) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.box_predictor.cls_score.bias' to the model due to incompatible shapes: (81,) in the checkpoint but (6,) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.box_predictor.bbox_pred.weight' to the model due to incompatible shapes: (320, 1024) in the checkpoint but (20, 1024) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.box_predictor.bbox_pred.bias' to the model due to incompatible shapes: (320,) in the checkpoint but (20,) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.mask_head.predictor.weight' to the model due to incompatible shapes: (80, 256, 1, 1) in the checkpoint but (5, 256, 1, 1) in the model! You might want to double check if this is expected.

Skip loading parameter 'roi_heads.mask_head.predictor.bias' to the model due to incompatible shapes: (80,) in the checkpoint but (5,) in the model! You might want to double check if this is expected.

Some model parameters or buffers are not found in the checkpoint:

`roi_heads.box_predictor.bbox_pred.{bias, weight}`

`roi_heads.box_predictor.cls_score.{bias, weight}`

`roi_heads.mask_head.predictor.{bias, weight}`

[05/19 16:17:44 d2.engine.train_loop]: Starting training from iteration 0
/usr/local/lib/python3.7/dist-packages/detectron2/structures/image_list.py:88: UserWarning: `__floordiv__` is deprecated, and its behavior will change in a future version of pytorch. It currently rounds toward 0 (like the 'trunc' function NOT 'floor'). This results in incorrect rounding for negative values. To keep the current behavior, use `torch.div(a, b, rounding_mode='trunc')`, or for actual floor division, use `torch.div(a, b, rounding_mode='floor')`.

`max_size = (max_size + (stride - 1)) // stride * stride`

[05/19 16:17:56 d2.utils.events]: eta: 0:22:58 iter: 19 total_loss: 2.912
loss_cls: 1.719 loss_box_reg: 0.388 loss_mask: 0.6955 loss_rpn_cls: 0.0530
1 loss_rpn_loc: 0.01793 time: 0.5663 data_time: 0.0148 lr: 2.1638e-06 max_mem: 1973M

[05/19 16:18:07 d2.utils.events]: eta: 0:22:35 iter: 39 total_loss: 2.94
loss_cls: 1.652 loss_box_reg: 0.4822 loss_mask: 0.6949 loss_rpn_cls: 0.0908
88 loss_rpn_loc: 0.02903 time: 0.5644 data_time: 0.0040 lr: 3.3888e-06 max_mem: 1973M

[05/19 16:18:19 d2.utils.events]: eta: 0:22:36 iter: 59 total_loss: 2.977
loss_cls: 1.574 loss_box_reg: 0.6284 loss_mask: 0.695 loss_rpn_cls: 0.0991
5 loss_rpn_loc: 0.02686 time: 0.5663 data_time: 0.0043 lr: 4.6137e-06 max_mem: 1973M

[05/19 16:18:30 d2.utils.events]: eta: 0:22:24 iter: 79 total_loss: 2.79
loss_cls: 1.462 loss_box_reg: 0.4704 loss_mask: 0.6932 loss_rpn_cls: 0.0720
08 loss_rpn_loc: 0.03234 time: 0.5629 data_time: 0.0040 lr: 5.8388e-06 max_mem: 1973M

[05/19 16:18:41 d2.utils.events]: eta: 0:22:08 iter: 99 total_loss: 2.596
loss_cls: 1.331 loss_box_reg: 0.4759 loss_mask: 0.6921 loss_rpn_cls: 0.0570
05 loss_rpn_loc: 0.02713 time: 0.5631 data_time: 0.0040 lr: 7.0638e-06 max_mem: 1973M

[05/19 16:18:53 d2.utils.events]: eta: 0:22:21 iter: 119 total_loss: 2.446
loss_cls: 1.199 loss_box_reg: 0.4031 loss_mask: 0.6905 loss_rpn_cls: 0.023

loss_rpn_loc: 0.0169 time: 0.5675 data_time: 0.0043 lr: 8.2888e-06 max_mem: 1973M

[05/19 16:19:04 d2.utils.events]: eta: 0:21:56 iter: 139 total_loss: 2.474
loss_cls: 1.094 loss_box_reg: 0.4957 loss_mask: 0.6878 loss_rpn_cls: 0.055
94 loss_rpn_loc: 0.02608 time: 0.5637 data_time: 0.0043 lr: 9.5138e-06 max_mem: 1973M

[05/19 16:19:15 d2.utils.events]: eta: 0:21:56 iter: 159 total_loss: 2.22
loss_cls: 0.9202 loss_box_reg: 0.4855 loss_mask: 0.6864 loss_rpn_cls: 0.05
018 loss_rpn_loc: 0.02128 time: 0.5659 data_time: 0.0043 lr: 1.0739e-05
max_mem: 1973M

[05/19 16:19:27 d2.utils.events]: eta: 0:21:49 iter: 179 total_loss: 2.061
loss_cls: 0.7976 loss_box_reg: 0.4309 loss_mask: 0.6854 loss_rpn_cls: 0.06
103 loss_rpn_loc: 0.02589 time: 0.5674 data_time: 0.0040 lr: 1.1964e-05
max_mem: 1973M

[05/19 16:19:38 d2.utils.events]: eta: 0:21:36 iter: 199 total_loss: 2.233
loss_cls: 0.8004 loss_box_reg: 0.6196 loss_mask: 0.6818 loss_rpn_cls: 0.05
38 loss_rpn_loc: 0.02512 time: 0.5668 data_time: 0.0040 lr: 1.3189e-05 max_mem: 1973M

/usr/local/lib/python3.7/dist-packages/fvcore/transforms/transform.py:724: ShapelyDeprecationWarning: Iteration over multi-part geometries is deprecated and will be removed in Shapely 2.0. Use the `geoms` property to access the constituent parts of a multi-part geometry.

for poly in cropped:

[05/19 16:19:49 d2.utils.events]: eta: 0:21:27 iter: 219 total_loss: 1.904
loss_cls: 0.6886 loss_box_reg: 0.4745 loss_mask: 0.6796 loss_rpn_cls: 0.06
429 loss_rpn_loc: 0.02604 time: 0.5651 data_time: 0.0042 lr: 1.4414e-05
max_mem: 1973M

[05/19 16:20:01 d2.utils.events]: eta: 0:21:16 iter: 239 total_loss: 1.795
loss_cls: 0.5902 loss_box_reg: 0.4523 loss_mask: 0.6688 loss_rpn_cls: 0.01
603 loss_rpn_loc: 0.01749 time: 0.5666 data_time: 0.0040 lr: 1.5639e-05
max_mem: 1973M

[05/19 16:20:12 d2.utils.events]: eta: 0:21:02 iter: 259 total_loss: 1.782
loss_cls: 0.5963 loss_box_reg: 0.4172 loss_mask: 0.669 loss_rpn_cls: 0.094
4 loss_rpn_loc: 0.03156 time: 0.5658 data_time: 0.0040 lr: 1.6864e-05 max_mem: 1973M

[05/19 16:20:23 d2.utils.events]: eta: 0:20:51 iter: 279 total_loss: 1.812
loss_cls: 0.6157 loss_box_reg: 0.4697 loss_mask: 0.6676 loss_rpn_cls: 0.02
427 loss_rpn_loc: 0.02499 time: 0.5669 data_time: 0.0043 lr: 1.8089e-05
max_mem: 1973M

[05/19 16:20:35 d2.utils.events]: eta: 0:20:36 iter: 299 total_loss: 1.829
loss_cls: 0.5827 loss_box_reg: 0.4743 loss_mask: 0.6683 loss_rpn_cls: 0.02
278 loss_rpn_loc: 0.01949 time: 0.5663 data_time: 0.0039 lr: 1.9314e-05
max_mem: 1973M

[05/19 16:20:47 d2.utils.events]: eta: 0:20:26 iter: 319 total_loss: 1.995
loss_cls: 0.6653 loss_box_reg: 0.6234 loss_mask: 0.6568 loss_rpn_cls: 0.02
385 loss_rpn_loc: 0.01612 time: 0.5680 data_time: 0.0039 lr: 2.0539e-05
max_mem: 1973M

[05/19 16:20:58 d2.utils.events]: eta: 0:20:14 iter: 339 total_loss: 1.596
loss_cls: 0.473 loss_box_reg: 0.4132 loss_mask: 0.6447 loss_rpn_cls: 0.028
39 loss_rpn_loc: 0.02253 time: 0.5677 data_time: 0.0041 lr: 2.1764e-05 max_mem: 1973M

[05/19 16:21:09 d2.utils.events]: eta: 0:20:03 iter: 359 total_loss: 1.708
loss_cls: 0.5335 loss_box_reg: 0.5021 loss_mask: 0.6519 loss_rpn_cls: 0.02

```

81  loss_rpn_loc: 0.02476  time: 0.5667  data_time: 0.0039  lr: 2.2989e-05  m
ax_mem: 1973M
[05/19 16:21:21 d2.utils.events]: eta: 0:19:54  iter: 379  total_loss: 1.868
loss_cls: 0.5813  loss_box_reg: 0.5318  loss_mask: 0.6363  loss_rpn_cls: 0.02
757  loss_rpn_loc: 0.02497  time: 0.5677  data_time: 0.0041  lr: 2.4214e-05
max_mem: 1973M
[05/19 16:21:34 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:21:34 d2.data.build]: Distribution of instances among all 5 categor
ies:
| category | #instances | category | #instances | category | #ins
tances |
|:-----|:-----|:-----|:-----|:-----|:-----|
| headlamp | 12          | rear_bumper | 5          | door     | 10
|
| hood      | 5           | front_bumper | 6          |          |
|
| total     | 38          |             |            |          |
|
[05/19 16:21:34 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:21:34 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...
[05/19 16:21:34 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:21:34 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:21:44 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0021 s/iter. Inference: 0.3789 s/iter. Eval: 0.5137 s/iter. Total: 0.8947 s
/iter. ETA=0:00:00
[05/19 16:21:44 d2.evaluation.evaluator]: Total inference time: 0:00:05.42828
1 (0.904714 s / iter per device, on 1 devices)
[05/19 16:21:44 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:02 (0.378862 s / iter per device, on 1 devices)
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Evaluating predictions with u
nofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:21:44 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:21:44 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.02 seconds.
[05/19 16:21:44 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:21:44 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.02 seconds.
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.0
07

```

```

Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.0
16
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.0
04
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.0
14
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.0
13
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.0
21
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.1
38
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.1
70
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.1
30
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.2
02
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP   | AP50 | AP75 | APs  | APm  | APl  |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 0.707 | 1.568 | 0.370 | nan  | 1.413 | 1.275 |
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category   | AP   | category   | AP   | category   | AP   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp   | 1.272 | rear_bumper | 1.052 | door        | 0.178 |
| hood        | 0.000 | front_bumper | 1.031 |              |      |
Loading and preparing results...
DONE (t=0.05s)
creating index...
index created!
[05/19 16:21:44 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[05/19 16:21:44 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.04 seconds.
[05/19 16:21:44 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:21:44 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.0
07
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.0
13
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.0
07
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000

```



```

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.0
06
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.0
16
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.0
47
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.1
32
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.1
60
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.1
40
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.1
87
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 0.726 | 1.253 | 0.731 | nan | 0.636 | 1.631 |
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:21:44 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 1.216 | rear_bumper | 0.704 | door | 0.068 |
| hood | 0.000 | front_bumper | 1.644 | | |
[05/19 16:21:44 d2.engine.defaults]: Evaluation results for car_part_val in c
sv format:
[05/19 16:21:44 d2.evaluation.testing]: cypypaste: Task: bbox
[05/19 16:21:44 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:21:44 d2.evaluation.testing]: cypypaste: 0.7067,1.5678,0.3703,nan,1
.4125,1.2746
[05/19 16:21:44 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:21:44 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:21:44 d2.evaluation.testing]: cypypaste: 0.7263,1.2533,0.7308,nan,0
.6364,1.6308
[05/19 16:21:44 d2.utils.events]: eta: 0:19:45 iter: 399 total_loss: 1.679
loss_cls: 0.4959 loss_box_reg: 0.5131 loss_mask: 0.6298 loss_rpn_cls: 0.01
686 loss_rpn_loc: 0.01959 time: 0.5687 data_time: 0.0040 lr: 2.5439e-05
max_mem: 1973M
[05/19 16:21:56 d2.utils.events]: eta: 0:19:34 iter: 419 total_loss: 2.006
loss_cls: 0.6237 loss_box_reg: 0.6319 loss_mask: 0.6455 loss_rpn_cls: 0.02
417 loss_rpn_loc: 0.01975 time: 0.5693 data_time: 0.0045 lr: 2.6664e-05
max_mem: 1973M
[05/19 16:22:07 d2.utils.events]: eta: 0:19:18 iter: 439 total_loss: 1.674
loss_cls: 0.5099 loss_box_reg: 0.5396 loss_mask: 0.628 loss_rpn_cls: 0.015
64 loss_rpn_loc: 0.01619 time: 0.5683 data_time: 0.0041 lr: 2.7889e-05 m
ax_mem: 1973M
[05/19 16:22:19 d2.utils.events]: eta: 0:19:11 iter: 459 total_loss: 1.879
loss_cls: 0.5805 loss_box_reg: 0.6008 loss_mask: 0.6191 loss_rpn_cls: 0.01
192 loss_rpn_loc: 0.01813 time: 0.5689 data_time: 0.0037 lr: 2.9114e-05
max_mem: 1973M

```

[05/19 16:22:30 d2.utils.events]: eta: 0:19:00 iter: 479 total_loss: 1.779
loss_cls: 0.5371 loss_box_reg: 0.4604 loss_mask: 0.6042 loss_rpn_cls: 0.02
69 loss_rpn_loc: 0.02171 time: 0.5690 data_time: 0.0040 lr: 3.0339e-05 m
ax_mem: 1973M

[05/19 16:22:41 d2.utils.events]: eta: 0:18:49 iter: 499 total_loss: 1.922
loss_cls: 0.6155 loss_box_reg: 0.7148 loss_mask: 0.6178 loss_rpn_cls: 0.01
176 loss_rpn_loc: 0.01917 time: 0.5690 data_time: 0.0040 lr: 3.1564e-05
max_mem: 1973M

[05/19 16:22:53 d2.utils.events]: eta: 0:18:33 iter: 519 total_loss: 1.703
loss_cls: 0.4749 loss_box_reg: 0.4438 loss_mask: 0.6062 loss_rpn_cls: 0.02
45 loss_rpn_loc: 0.01799 time: 0.5685 data_time: 0.0037 lr: 3.2789e-05 m
ax_mem: 1973M

[05/19 16:23:04 d2.utils.events]: eta: 0:18:21 iter: 539 total_loss: 1.77
loss_cls: 0.5429 loss_box_reg: 0.5844 loss_mask: 0.6112 loss_rpn_cls: 0.02
495 loss_rpn_loc: 0.02221 time: 0.5682 data_time: 0.0040 lr: 3.4014e-05
max_mem: 1973M

[05/19 16:23:15 d2.utils.events]: eta: 0:18:10 iter: 559 total_loss: 1.538
loss_cls: 0.4654 loss_box_reg: 0.4335 loss_mask: 0.5902 loss_rpn_cls: 0.01
866 loss_rpn_loc: 0.0174 time: 0.5681 data_time: 0.0037 lr: 3.5239e-05 m
ax_mem: 1973M

[05/19 16:23:27 d2.utils.events]: eta: 0:17:59 iter: 579 total_loss: 1.893
loss_cls: 0.6398 loss_box_reg: 0.6764 loss_mask: 0.576 loss_rpn_cls: 0.023
53 loss_rpn_loc: 0.01981 time: 0.5681 data_time: 0.0041 lr: 3.6464e-05 m
ax_mem: 1973M

[05/19 16:23:38 d2.utils.events]: eta: 0:17:48 iter: 599 total_loss: 1.738
loss_cls: 0.5284 loss_box_reg: 0.6531 loss_mask: 0.5683 loss_rpn_cls: 0.00
952 loss_rpn_loc: 0.01653 time: 0.5687 data_time: 0.0040 lr: 3.7689e-05
max_mem: 1973M

[05/19 16:23:50 d2.utils.events]: eta: 0:17:37 iter: 619 total_loss: 1.78
loss_cls: 0.5284 loss_box_reg: 0.6026 loss_mask: 0.5701 loss_rpn_cls: 0.03
085 loss_rpn_loc: 0.02204 time: 0.5693 data_time: 0.0039 lr: 3.8914e-05
max_mem: 1973M

[05/19 16:24:01 d2.utils.events]: eta: 0:17:26 iter: 639 total_loss: 1.615
loss_cls: 0.4581 loss_box_reg: 0.5543 loss_mask: 0.5523 loss_rpn_cls: 0.01
749 loss_rpn_loc: 0.01201 time: 0.5694 data_time: 0.0040 lr: 4.0139e-05
max_mem: 1973M

[05/19 16:24:13 d2.utils.events]: eta: 0:17:17 iter: 659 total_loss: 1.901
loss_cls: 0.5834 loss_box_reg: 0.6789 loss_mask: 0.5408 loss_rpn_cls: 0.01
509 loss_rpn_loc: 0.02078 time: 0.5700 data_time: 0.0040 lr: 4.1364e-05
max_mem: 1973M

[05/19 16:24:25 d2.utils.events]: eta: 0:17:03 iter: 679 total_loss: 1.662
loss_cls: 0.5019 loss_box_reg: 0.56 loss_mask: 0.5771 loss_rpn_cls: 0.0250
8 loss_rpn_loc: 0.01921 time: 0.5699 data_time: 0.0039 lr: 4.2589e-05 m
x_mem: 1973M

[05/19 16:24:36 d2.utils.events]: eta: 0:16:51 iter: 699 total_loss: 1.684
loss_cls: 0.5184 loss_box_reg: 0.5737 loss_mask: 0.518 loss_rpn_cls: 0.013
77 loss_rpn_loc: 0.01477 time: 0.5695 data_time: 0.0044 lr: 4.3814e-05 m
ax_mem: 1973M

[05/19 16:24:47 d2.utils.events]: eta: 0:16:40 iter: 719 total_loss: 1.817
loss_cls: 0.5514 loss_box_reg: 0.7039 loss_mask: 0.5124 loss_rpn_cls: 0.01
088 loss_rpn_loc: 0.01296 time: 0.5696 data_time: 0.0042 lr: 4.5039e-05
max_mem: 1973M

```

[05/19 16:24:59 d2.utils.events]: eta: 0:16:30 iter: 739 total_loss: 1.578
loss_cls: 0.4883 loss_box_reg: 0.5709 loss_mask: 0.4762 loss_rpn_cls: 0.01
497 loss_rpn_loc: 0.01215 time: 0.5702 data_time: 0.0039 lr: 4.6264e-05
max_mem: 1973M
[05/19 16:25:11 d2.utils.events]: eta: 0:16:21 iter: 759 total_loss: 1.788
loss_cls: 0.5291 loss_box_reg: 0.6949 loss_mask: 0.5273 loss_rpn_cls: 0.02
248 loss_rpn_loc: 0.02161 time: 0.5706 data_time: 0.0041 lr: 4.7489e-05
max_mem: 1973M
[05/19 16:25:22 d2.utils.events]: eta: 0:16:08 iter: 779 total_loss: 1.75
loss_cls: 0.5302 loss_box_reg: 0.6109 loss_mask: 0.5269 loss_rpn_cls: 0.01
418 loss_rpn_loc: 0.01828 time: 0.5702 data_time: 0.0040 lr: 4.8714e-05
max_mem: 1973M
[05/19 16:25:35 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:25:35 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:25:35 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...
[05/19 16:25:35 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:25:35 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:25:45 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0017 s/iter. Inference: 0.3804 s/iter. Eval: 0.4693 s/iter. Total: 0.8513 s
/iter. ETA=0:00:00
[05/19 16:25:45 d2.evaluation.evaluator]: Total inference time: 0:00:05.22639
6 (0.871066 s / iter per device, on 1 devices)
[05/19 16:25:45 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:02 (0.380384 s / iter per device, on 1 devices)
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Evaluating predictions with u
nofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:25:45 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:25:45 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.04 seconds.
[05/19 16:25:45 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:25:45 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.05 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.0
59
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.1
19
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.0
46
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.2
16
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.0
80
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.0
74
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.2
99
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.3
57
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.3
67
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.3
77
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 5.898 | 11.873 | 4.612 | nan | 21.611 | 8.010 |
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 13.549 | rear_bumper | 8.484 | door | 2.181 |
| hood | 0.861 | front_bumper | 4.416 | | |
Loading and preparing results...
DONE (t=0.07s)
creating index...
index created!
[05/19 16:25:45 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[05/19 16:25:45 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.06 seconds.
[05/19 16:25:45 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:25:45 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.03 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.0
62
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.1
02
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.0
70
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.0
94
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.0
86
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.0
79

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.2
76
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.3
42
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.3
55
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.3
62
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP   | AP50  | AP75  | APs   | APm   | APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 6.158 | 10.192 | 7.000 | nan   | 9.412 | 8.574 |
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:25:45 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP      | category | AP      | category | AP      |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 14.620 | rear_bumper | 9.009 | door      | 1.361 |
| hood     | 0.598  | front_bumper | 5.201 |           |       |
[05/19 16:25:45 d2.engine.defaults]: Evaluation results for car_part_val in c
sv format:
[05/19 16:25:45 d2.evaluation.testing]: cypypaste: Task: bbox
[05/19 16:25:45 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:25:45 d2.evaluation.testing]: cypypaste: 5.8981,11.8728,4.6115,nan,
21.6106,8.0098
[05/19 16:25:45 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:25:45 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:25:45 d2.evaluation.testing]: cypypaste: 6.1579,10.1923,6.9999,nan,
9.4122,8.5737
[05/19 16:25:45 d2.utils.events]: eta: 0:15:58 iter: 799 total_loss: 1.654
loss_cls: 0.5052 loss_box_reg: 0.5849 loss_mask: 0.471 loss_rpn_cls: 0.013
37 loss_rpn_loc: 0.01942 time: 0.5707 data_time: 0.0039 lr: 4.9939e-05 m
ax_mem: 1973M
[05/19 16:25:57 d2.utils.events]: eta: 0:15:48 iter: 819 total_loss: 1.576
loss_cls: 0.4259 loss_box_reg: 0.5673 loss_mask: 0.4915 loss_rpn_cls: 0.01
456 loss_rpn_loc: 0.01705 time: 0.5710 data_time: 0.0044 lr: 5e-05 max_m
em: 1973M
[05/19 16:26:09 d2.utils.events]: eta: 0:15:36 iter: 839 total_loss: 1.687
loss_cls: 0.4692 loss_box_reg: 0.5968 loss_mask: 0.4919 loss_rpn_cls: 0.01
449 loss_rpn_loc: 0.01868 time: 0.5713 data_time: 0.0041 lr: 5e-05 max_m
em: 1973M
[05/19 16:26:21 d2.utils.events]: eta: 0:15:27 iter: 859 total_loss: 1.636
loss_cls: 0.4951 loss_box_reg: 0.5674 loss_mask: 0.476 loss_rpn_cls: 0.011
76 loss_rpn_loc: 0.01603 time: 0.5719 data_time: 0.0043 lr: 5e-05 max_me
m: 1973M
[05/19 16:26:32 d2.utils.events]: eta: 0:15:16 iter: 879 total_loss: 1.632
loss_cls: 0.4648 loss_box_reg: 0.6361 loss_mask: 0.4416 loss_rpn_cls: 0.01
078 loss_rpn_loc: 0.01448 time: 0.5721 data_time: 0.0037 lr: 5e-05 max_m
em: 1973M
[05/19 16:26:44 d2.utils.events]: eta: 0:15:04 iter: 899 total_loss: 1.336
loss_cls: 0.3644 loss_box_reg: 0.5517 loss_mask: 0.4835 loss_rpn_cls: 0.00

```

8206 loss_rpn_loc: 0.01477 time: 0.5720 data_time: 0.0038 lr: 5e-05 max_mem: 1973M
[05/19 16:26:55 d2.utils.events]: eta: 0:14:53 iter: 919 total_loss: 1.624
loss_cls: 0.4793 loss_box_reg: 0.6665 loss_mask: 0.4269 loss_rpn_cls: 0.01
634 loss_rpn_loc: 0.0154 time: 0.5720 data_time: 0.0040 lr: 5e-05 max_m
m: 1973M
[05/19 16:27:07 d2.utils.events]: eta: 0:14:42 iter: 939 total_loss: 1.511
loss_cls: 0.453 loss_box_reg: 0.6032 loss_mask: 0.4347 loss_rpn_cls: 0.008
292 loss_rpn_loc: 0.01305 time: 0.5725 data_time: 0.0041 lr: 5e-05 max_m
em: 1973M
[05/19 16:27:19 d2.utils.events]: eta: 0:14:31 iter: 959 total_loss: 1.508
loss_cls: 0.4651 loss_box_reg: 0.6068 loss_mask: 0.4207 loss_rpn_cls: 0.00
7813 loss_rpn_loc: 0.013 time: 0.5729 data_time: 0.0040 lr: 5e-05 max_me
m: 1973M
[05/19 16:27:31 d2.utils.events]: eta: 0:14:21 iter: 979 total_loss: 1.63
loss_cls: 0.4982 loss_box_reg: 0.7252 loss_mask: 0.3903 loss_rpn_cls: 0.00
5221 loss_rpn_loc: 0.01375 time: 0.5732 data_time: 0.0040 lr: 5e-05 max_
mem: 1973M
[05/19 16:27:42 d2.utils.events]: eta: 0:14:09 iter: 999 total_loss: 1.447
loss_cls: 0.4453 loss_box_reg: 0.5296 loss_mask: 0.4043 loss_rpn_cls: 0.01
307 loss_rpn_loc: 0.01279 time: 0.5734 data_time: 0.0038 lr: 5e-05 max_m
em: 1973M
[05/19 16:27:54 d2.utils.events]: eta: 0:13:58 iter: 1019 total_loss: 1.53
4 loss_cls: 0.4484 loss_box_reg: 0.6541 loss_mask: 0.4062 loss_rpn_cls: 0
.01049 loss_rpn_loc: 0.01563 time: 0.5736 data_time: 0.0038 lr: 5e-05 ma
x_mem: 1973M
[05/19 16:28:06 d2.utils.events]: eta: 0:13:48 iter: 1039 total_loss: 1.56
3 loss_cls: 0.4441 loss_box_reg: 0.6367 loss_mask: 0.377 loss_rpn_cls: 0.
01439 loss_rpn_loc: 0.02127 time: 0.5742 data_time: 0.0037 lr: 5e-05 max
_mem: 1973M
[05/19 16:28:18 d2.utils.events]: eta: 0:13:37 iter: 1059 total_loss: 1.56
8 loss_cls: 0.4703 loss_box_reg: 0.6634 loss_mask: 0.3592 loss_rpn_cls: 0
.007671 loss_rpn_loc: 0.01275 time: 0.5742 data_time: 0.0048 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:28:29 d2.utils.events]: eta: 0:13:26 iter: 1079 total_loss: 1.39
4 loss_cls: 0.3938 loss_box_reg: 0.6155 loss_mask: 0.3718 loss_rpn_cls: 0
.01028 loss_rpn_loc: 0.01061 time: 0.5744 data_time: 0.0039 lr: 5e-05 ma
x_mem: 1973M
[05/19 16:28:41 d2.utils.events]: eta: 0:13:15 iter: 1099 total_loss: 1.45
5 loss_cls: 0.3879 loss_box_reg: 0.6113 loss_mask: 0.3384 loss_rpn_cls: 0
.00731 loss_rpn_loc: 0.01341 time: 0.5744 data_time: 0.0037 lr: 5e-05 ma
x_mem: 1973M
[05/19 16:28:52 d2.utils.events]: eta: 0:13:03 iter: 1119 total_loss: 1.45
4 loss_cls: 0.4259 loss_box_reg: 0.5707 loss_mask: 0.3253 loss_rpn_cls: 0
.01511 loss_rpn_loc: 0.02098 time: 0.5744 data_time: 0.0040 lr: 5e-05 ma
x_mem: 1973M
[05/19 16:29:04 d2.utils.events]: eta: 0:12:53 iter: 1139 total_loss: 1.39
loss_cls: 0.3677 loss_box_reg: 0.6058 loss_mask: 0.315 loss_rpn_cls: 0.009
18 loss_rpn_loc: 0.01608 time: 0.5747 data_time: 0.0038 lr: 5e-05 max_me
m: 1973M
[05/19 16:29:16 d2.utils.events]: eta: 0:12:41 iter: 1159 total_loss: 1.34
4 loss_cls: 0.3841 loss_box_reg: 0.5768 loss_mask: 0.343 loss_rpn_cls: 0.


```

009491 loss_rpn_loc: 0.0191 time: 0.5750 data_time: 0.0039 lr: 5e-05 max
_mem: 1973M
[05/19 16:29:28 d2.utils.events]: eta: 0:12:30 iter: 1179 total_loss: 1.37
9 loss_cls: 0.3609 loss_box_reg: 0.6155 loss_mask: 0.3188 loss_rpn_cls: 0
.01023 loss_rpn_loc: 0.01353 time: 0.5750 data_time: 0.0041 lr: 5e-05 ma
x_mem: 1973M
[05/19 16:29:40 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:29:40 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:29:40 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...
[05/19 16:29:40 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:29:40 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:29:48 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0018 s/iter. Inference: 0.3441 s/iter. Eval: 0.3631 s/iter. Total: 0.7090 s
/iter. ETA=0:00:00
[05/19 16:29:49 d2.evaluation.evaluator]: Total inference time: 0:00:04.31249
3 (0.718749 s / iter per device, on 1 devices)
[05/19 16:29:49 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:02 (0.344144 s / iter per device, on 1 devices)
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Evaluating predictions with u
nofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:29:49 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:29:49 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.02 seconds.
[05/19 16:29:49 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:29:49 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.1
10
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.2
52
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.0
77
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.3
10
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.1
27
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.1
48

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.4
17
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.4
37
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.3
32
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.4
98
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
|  AP   | AP50 | AP75 | APs  | APm   | APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 10.983 | 25.187 | 7.696 | nan  | 30.970 | 12.670 |
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP      | category | AP      | category | AP      |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 18.129 | rear_bumper | 8.686 | door      | 11.230 |
| hood      | 3.993 | front_bumper | 12.877 |           |         |
Loading and preparing results...
DONE (t=0.02s)
creating index...
index created!
[05/19 16:29:49 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[05/19 16:29:49 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.03 seconds.
[05/19 16:29:49 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:29:49 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.02 seconds.
Average Precision   (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.1
67
Average Precision   (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.2
62
Average Precision   (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.2
04
Average Precision   (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision   (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.3
99
Average Precision   (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.1
84
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.2
71
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.4
74
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.4
87
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
46
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.5
17
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Evaluation results for segm:
|  AP   |  AP50  |  AP75  |  APs   |  APm   |  APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 16.742 | 26.240 | 20.368 | nan    | 39.924 | 18.378 |
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:29:49 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP      | category | AP      | category | AP      |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 21.344 | rear_bumper | 11.924 | door      | 18.698 |
| hood      | 9.675  | front_bumper | 22.069 |           |         |
[05/19 16:29:49 d2.engine.defaults]: Evaluation results for car_part_val in c
sv format:
[05/19 16:29:49 d2.evaluation.testing]: cypypaste: Task: bbox
[05/19 16:29:49 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:29:49 d2.evaluation.testing]: cypypaste: 10.9831,25.1871,7.6958,nan
,30.9703,12.6704
[05/19 16:29:49 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:29:49 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:29:49 d2.evaluation.testing]: cypypaste: 16.7420,26.2404,20.3683,na
n,39.9236,18.3775
[05/19 16:29:49 d2.utils.events]: eta: 0:12:18 iter: 1199 total_loss: 1.35
8 loss_cls: 0.3763 loss_box_reg: 0.5969 loss_mask: 0.2697 loss_rpn_cls: 0
.01445 loss_rpn_loc: 0.0133 time: 0.5750 data_time: 0.0038 lr: 5e-05 max
_mem: 1973M
[05/19 16:30:01 d2.utils.events]: eta: 0:12:08 iter: 1219 total_loss: 1.39
2 loss_cls: 0.3834 loss_box_reg: 0.6443 loss_mask: 0.2988 loss_rpn_cls: 0
.01985 loss_rpn_loc: 0.0182 time: 0.5754 data_time: 0.0041 lr: 5e-05 max
_mem: 1973M
[05/19 16:30:12 d2.utils.events]: eta: 0:11:57 iter: 1239 total_loss: 1.19
3 loss_cls: 0.3186 loss_box_reg: 0.5825 loss_mask: 0.2567 loss_rpn_cls: 0
.007344 loss_rpn_loc: 0.01505 time: 0.5752 data_time: 0.0037 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:30:24 d2.utils.events]: eta: 0:11:46 iter: 1259 total_loss: 1.35
2 loss_cls: 0.3608 loss_box_reg: 0.6159 loss_mask: 0.2717 loss_rpn_cls: 0
.01479 loss_rpn_loc: 0.0185 time: 0.5754 data_time: 0.0039 lr: 5e-05 max
_mem: 1973M
[05/19 16:30:35 d2.utils.events]: eta: 0:11:34 iter: 1279 total_loss: 1.24
7 loss_cls: 0.3232 loss_box_reg: 0.5642 loss_mask: 0.2664 loss_rpn_cls: 0
.004215 loss_rpn_loc: 0.01309 time: 0.5754 data_time: 0.0043 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:30:48 d2.utils.events]: eta: 0:11:25 iter: 1299 total_loss: 1.16
3 loss_cls: 0.3041 loss_box_reg: 0.5601 loss_mask: 0.2539 loss_rpn_cls: 0
.007191 loss_rpn_loc: 0.01105 time: 0.5758 data_time: 0.0038 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:30:59 d2.utils.events]: eta: 0:11:15 iter: 1319 total_loss: 1.24
6 loss_cls: 0.3141 loss_box_reg: 0.552 loss_mask: 0.2717 loss_rpn_cls: 0.
01647 loss_rpn_loc: 0.02313 time: 0.5761 data_time: 0.0041 lr: 5e-05 max
_mem: 1973M

```

[05/19 16:31:11 d2.utils.events]: eta: 0:11:04 iter: 1339 total_loss: 1.07
6 loss_cls: 0.2722 loss_box_reg: 0.5406 loss_mask: 0.2461 loss_rpn_cls: 0
.003541 loss_rpn_loc: 0.01514 time: 0.5763 data_time: 0.0040 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:31:23 d2.utils.events]: eta: 0:10:52 iter: 1359 total_loss: 1.14
7 loss_cls: 0.3112 loss_box_reg: 0.5273 loss_mask: 0.2399 loss_rpn_cls: 0
.002606 loss_rpn_loc: 0.01359 time: 0.5763 data_time: 0.0038 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:31:35 d2.utils.events]: eta: 0:10:41 iter: 1379 total_loss: 1.05
8 loss_cls: 0.2627 loss_box_reg: 0.5083 loss_mask: 0.2268 loss_rpn_cls: 0
.003226 loss_rpn_loc: 0.01657 time: 0.5765 data_time: 0.0037 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:31:46 d2.utils.events]: eta: 0:10:30 iter: 1399 total_loss: 1.15
6 loss_cls: 0.3071 loss_box_reg: 0.5028 loss_mask: 0.2443 loss_rpn_cls: 0
.006123 loss_rpn_loc: 0.0191 time: 0.5767 data_time: 0.0038 lr: 5e-05 ma
x_mem: 1973M

[05/19 16:31:58 d2.utils.events]: eta: 0:10:18 iter: 1419 total_loss: 1.02
7 loss_cls: 0.3091 loss_box_reg: 0.4665 loss_mask: 0.2057 loss_rpn_cls: 0
.005477 loss_rpn_loc: 0.01548 time: 0.5767 data_time: 0.0039 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:32:10 d2.utils.events]: eta: 0:10:08 iter: 1439 total_loss: 0.91
15 loss_cls: 0.2786 loss_box_reg: 0.4367 loss_mask: 0.2042 loss_rpn_cls:
0.003848 loss_rpn_loc: 0.02047 time: 0.5773 data_time: 0.0043 lr: 5e-05
max_mem: 1973M

[05/19 16:32:22 d2.utils.events]: eta: 0:09:56 iter: 1459 total_loss: 1.06
6 loss_cls: 0.2925 loss_box_reg: 0.4858 loss_mask: 0.2318 loss_rpn_cls: 0
.01086 loss_rpn_loc: 0.01477 time: 0.5772 data_time: 0.0039 lr: 5e-05 ma
x_mem: 1973M

[05/19 16:32:34 d2.utils.events]: eta: 0:09:45 iter: 1479 total_loss: 0.81
48 loss_cls: 0.234 loss_box_reg: 0.3763 loss_mask: 0.1899 loss_rpn_cls: 0
.002402 loss_rpn_loc: 0.01685 time: 0.5774 data_time: 0.0039 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:32:46 d2.utils.events]: eta: 0:09:34 iter: 1499 total_loss: 1.07
7 loss_cls: 0.3001 loss_box_reg: 0.4956 loss_mask: 0.2005 loss_rpn_cls: 0
.009581 loss_rpn_loc: 0.02184 time: 0.5776 data_time: 0.0037 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:32:57 d2.utils.events]: eta: 0:09:23 iter: 1519 total_loss: 0.88
64 loss_cls: 0.2222 loss_box_reg: 0.4327 loss_mask: 0.2237 loss_rpn_cls:
0.00247 loss_rpn_loc: 0.01417 time: 0.5774 data_time: 0.0039 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:33:08 d2.utils.events]: eta: 0:09:11 iter: 1539 total_loss: 0.85
47 loss_cls: 0.2173 loss_box_reg: 0.4033 loss_mask: 0.1721 loss_rpn_cls:
0.005746 loss_rpn_loc: 0.01305 time: 0.5774 data_time: 0.0038 lr: 5e-05
max_mem: 1973M

[05/19 16:33:20 d2.utils.events]: eta: 0:09:00 iter: 1559 total_loss: 0.87
77 loss_cls: 0.2639 loss_box_reg: 0.3845 loss_mask: 0.1809 loss_rpn_cls:
0.004612 loss_rpn_loc: 0.0131 time: 0.5773 data_time: 0.0042 lr: 5e-05 m
ax_mem: 1973M

[05/19 16:33:31 d2.utils.events]: eta: 0:08:48 iter: 1579 total_loss: 0.88
54 loss_cls: 0.2309 loss_box_reg: 0.3992 loss_mask: 0.2042 loss_rpn_cls:
0.003125 loss_rpn_loc: 0.01559 time: 0.5773 data_time: 0.0038 lr: 5e-05
max_mem: 1973M

```

[05/19 16:33:44 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:33:44 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:33:44 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...
[05/19 16:33:44 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:33:44 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:33:51 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0018 s/iter. Inference: 0.3182 s/iter. Eval: 0.1946 s/iter. Total: 0.5146 s
/iter. ETA=0:00:00
[05/19 16:33:51 d2.evaluation.evaluator]: Total inference time: 0:00:03.17825
2 (0.529709 s / iter per device, on 1 devices)
[05/19 16:33:51 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:01 (0.318199 s / iter per device, on 1 devices)
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Evaluating predictions with u
nofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:33:51 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:33:51 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.03 seconds.
[05/19 16:33:51 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:33:51 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.04 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.2
91
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.5
04
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.2
57
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
22
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.2
68
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.3
68
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.5
23
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.5
26
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
57
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.5
56
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
|  AP   |  AP50  |  AP75  |  APs   |  APm   |  APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 29.125 | 50.371 | 25.720 | nan    | 42.205 | 26.806 |
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP      | category | AP      | category | AP      |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 31.565 | rear_bumper | 22.773 | door      | 33.574 |
| hood      | 33.011 | front_bumper | 24.702 |           |         |
Loading and preparing results...
DONE (t=0.02s)
creating index...
index created!
[05/19 16:33:51 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[05/19 16:33:51 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.03 seconds.
[05/19 16:33:51 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:33:51 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.03 seconds.
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.3
40
Average Precision  (AP) @[ IoU=0.50      | area=  all | maxDets=100 ] = 0.4
82
Average Precision  (AP) @[ IoU=0.75      | area=  all | maxDets=100 ] = 0.3
86
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
84
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.3
00
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets= 1 ] = 0.4
11
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets= 10 ] = 0.5
22
Average Recall     (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.5
25
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.5
40
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.5
24
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Evaluation results for segm:
|  AP   |  AP50  |  AP75  |  APs   |  APm   |  APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|

```

```

| 34.010 | 48.189 | 38.610 | nan | 48.440 | 29.979 |
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/19 16:33:51 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP | category | AP | category | AP |
|:-----|:-----|:-----|:-----|:-----|:-----|
| headlamp | 30.611 | rear_bumper | 21.404 | door | 43.500 |
| hood | 36.845 | front_bumper | 37.690 | | |
[05/19 16:33:51 d2.engine.defaults]: Evaluation results for car_part_val in csv format:
[05/19 16:33:51 d2.evaluation.testing]: cypypaste: Task: bbox
[05/19 16:33:51 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:33:51 d2.evaluation.testing]: cypypaste: 29.1247,50.3714,25.7201,nan,42.2046,26.8057
[05/19 16:33:51 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:33:51 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:33:51 d2.evaluation.testing]: cypypaste: 34.0100,48.1887,38.6096,nan,48.4404,29.9786
[05/19 16:33:51 d2.utils.events]: eta: 0:08:37 iter: 1599 total_loss: 0.8612 loss_cls: 0.2212 loss_box_reg: 0.3836 loss_mask: 0.1931 loss_rpn_cls: 0.003225 loss_rpn_loc: 0.01717 time: 0.5774 data_time: 0.0038 lr: 5e-05 max_mem: 1973M
[05/19 16:34:03 d2.utils.events]: eta: 0:08:26 iter: 1619 total_loss: 0.7946 loss_cls: 0.1946 loss_box_reg: 0.3684 loss_mask: 0.1845 loss_rpn_cls: 0.003005 loss_rpn_loc: 0.01859 time: 0.5779 data_time: 0.0043 lr: 5e-05 max_mem: 1973M
[05/19 16:34:15 d2.utils.events]: eta: 0:08:14 iter: 1639 total_loss: 0.7232 loss_cls: 0.2119 loss_box_reg: 0.3667 loss_mask: 0.1762 loss_rpn_cls: 0.002419 loss_rpn_loc: 0.01465 time: 0.5777 data_time: 0.0041 lr: 5e-05 max_mem: 1973M
[05/19 16:34:27 d2.utils.events]: eta: 0:08:03 iter: 1659 total_loss: 0.7122 loss_cls: 0.2015 loss_box_reg: 0.3879 loss_mask: 0.1729 loss_rpn_cls: 0.007644 loss_rpn_loc: 0.01507 time: 0.5779 data_time: 0.0037 lr: 5e-05 max_mem: 1973M
[05/19 16:34:38 d2.utils.events]: eta: 0:07:53 iter: 1679 total_loss: 0.7054 loss_cls: 0.1775 loss_box_reg: 0.3209 loss_mask: 0.1508 loss_rpn_cls: 0.004641 loss_rpn_loc: 0.01073 time: 0.5780 data_time: 0.0039 lr: 5e-05 max_mem: 1973M
[05/19 16:34:50 d2.utils.events]: eta: 0:07:41 iter: 1699 total_loss: 0.7431 loss_cls: 0.181 loss_box_reg: 0.3728 loss_mask: 0.1941 loss_rpn_cls: 0.006242 loss_rpn_loc: 0.0174 time: 0.5779 data_time: 0.0039 lr: 5e-05 max_mem: 1973M
[05/19 16:35:01 d2.utils.events]: eta: 0:07:30 iter: 1719 total_loss: 0.8182 loss_cls: 0.2189 loss_box_reg: 0.3448 loss_mask: 0.1693 loss_rpn_cls: 0.01335 loss_rpn_loc: 0.0181 time: 0.5776 data_time: 0.0038 lr: 5e-05 max_mem: 1973M
[05/19 16:35:13 d2.utils.events]: eta: 0:07:18 iter: 1739 total_loss: 0.7425 loss_cls: 0.179 loss_box_reg: 0.3365 loss_mask: 0.1874 loss_rpn_cls: 0.001968 loss_rpn_loc: 0.01409 time: 0.5779 data_time: 0.0042 lr: 5e-05 max_mem: 1973M
[05/19 16:35:25 d2.utils.events]: eta: 0:07:06 iter: 1759 total_loss: 0.7746 loss_cls: 0.1957 loss_box_reg: 0.3579 loss_mask: 0.1805 loss_rpn_cls:

```

0.001925 loss_rpn_loc: 0.01719 time: 0.5783 data_time: 0.0038 lr: 5e-05
max_mem: 1973M
[05/19 16:35:36 d2.utils.events]: eta: 0:06:55 iter: 1779 total_loss: 0.68
3 loss_cls: 0.1724 loss_box_reg: 0.3378 loss_mask: 0.1588 loss_rpn_cls: 0
.005584 loss_rpn_loc: 0.01703 time: 0.5780 data_time: 0.0037 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:35:48 d2.utils.events]: eta: 0:06:43 iter: 1799 total_loss: 0.68
04 loss_cls: 0.1594 loss_box_reg: 0.3313 loss_mask: 0.1376 loss_rpn_cls:
0.003003 loss_rpn_loc: 0.0133 time: 0.5781 data_time: 0.0039 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:35:59 d2.utils.events]: eta: 0:06:31 iter: 1819 total_loss: 0.80
21 loss_cls: 0.1872 loss_box_reg: 0.3782 loss_mask: 0.1791 loss_rpn_cls:
0.006496 loss_rpn_loc: 0.0139 time: 0.5779 data_time: 0.0040 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:36:11 d2.utils.events]: eta: 0:06:20 iter: 1839 total_loss: 0.60
35 loss_cls: 0.1458 loss_box_reg: 0.2211 loss_mask: 0.1479 loss_rpn_cls:
0.001613 loss_rpn_loc: 0.009983 time: 0.5780 data_time: 0.0039 lr: 5e-05
max_mem: 1973M
[05/19 16:36:23 d2.utils.events]: eta: 0:06:08 iter: 1859 total_loss: 0.68
76 loss_cls: 0.1643 loss_box_reg: 0.3071 loss_mask: 0.1661 loss_rpn_cls:
0.002917 loss_rpn_loc: 0.01456 time: 0.5781 data_time: 0.0038 lr: 5e-05
max_mem: 1973M
[05/19 16:36:34 d2.utils.events]: eta: 0:05:57 iter: 1879 total_loss: 0.71
87 loss_cls: 0.1678 loss_box_reg: 0.3363 loss_mask: 0.1622 loss_rpn_cls:
0.0009991 loss_rpn_loc: 0.01448 time: 0.5781 data_time: 0.0039 lr: 5e-05
max_mem: 1973M
[05/19 16:36:46 d2.utils.events]: eta: 0:05:46 iter: 1899 total_loss: 0.65
89 loss_cls: 0.1693 loss_box_reg: 0.2975 loss_mask: 0.1414 loss_rpn_cls:
0.003699 loss_rpn_loc: 0.01363 time: 0.5782 data_time: 0.0039 lr: 5e-05
max_mem: 1973M
[05/19 16:36:58 d2.utils.events]: eta: 0:05:35 iter: 1919 total_loss: 0.69
59 loss_cls: 0.1645 loss_box_reg: 0.325 loss_mask: 0.1622 loss_rpn_cls: 0
.007974 loss_rpn_loc: 0.01651 time: 0.5784 data_time: 0.0040 lr: 5e-05 m
ax_mem: 1973M
[05/19 16:37:10 d2.utils.events]: eta: 0:05:23 iter: 1939 total_loss: 0.73
21 loss_cls: 0.1723 loss_box_reg: 0.3294 loss_mask: 0.2008 loss_rpn_cls:
0.009495 loss_rpn_loc: 0.01195 time: 0.5786 data_time: 0.0042 lr: 5e-05
max_mem: 1973M
[05/19 16:37:22 d2.utils.events]: eta: 0:05:12 iter: 1959 total_loss: 0.65
65 loss_cls: 0.1797 loss_box_reg: 0.2884 loss_mask: 0.1473 loss_rpn_cls:
0.002564 loss_rpn_loc: 0.01295 time: 0.5786 data_time: 0.0044 lr: 5e-05
max_mem: 1973M
[05/19 16:37:33 d2.utils.events]: eta: 0:05:00 iter: 1979 total_loss: 0.59
47 loss_cls: 0.1529 loss_box_reg: 0.2673 loss_mask: 0.1379 loss_rpn_cls:
0.002535 loss_rpn_loc: 0.01477 time: 0.5785 data_time: 0.0040 lr: 5e-05
max_mem: 1973M
[05/19 16:37:46 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:37:46 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:37:46 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...


```

[05/19 16:37:46 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:37:46 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:37:51 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0016 s/iter. Inference: 0.3012 s/iter. Eval: 0.1513 s/iter. Total: 0.4541 s
/iter. ETA=0:00:00
[05/19 16:37:51 d2.evaluation.evaluator]: Total inference time: 0:00:02.78967
1 (0.464945 s / iter per device, on 1 devices)
[05/19 16:37:51 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:01 (0.301229 s / iter per device, on 1 devices)
[05/19 16:37:51 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:37:51 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
[05/19 16:37:51 d2.evaluation.coco_evaluation]: Evaluating predictions with u
nofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:37:51 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:37:51 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.02 seconds.
[05/19 16:37:51 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:37:51 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.3
09
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.5
29
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.2
63
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
16
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.2
79
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.3
48
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.5
08
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.5
15
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
82
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.5
29
[05/19 16:37:51 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|

```

```

| 30.948 | 52.880 | 26.270 | nan | 41.602 | 27.902 |
[05/19 16:37:51 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/19 16:37:51 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----|:-----|:-----|:-----|:-----|:-----|
| headlamp | 35.509 | rear_bumper | 8.583 | door | 46.110 |
| hood | 39.797 | front_bumper | 24.738 | | |
Loading and preparing results...
DONE (t=0.01s)
creating index...
index created!
[05/19 16:37:52 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[05/19 16:37:52 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.02 seconds.
[05/19 16:37:52 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:37:52 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.346
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.508
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.386
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.504
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.309
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.494
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.496
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.544
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.500
[05/19 16:37:52 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----|:-----|:-----|:-----|:-----|:-----|
| 34.630 | 50.803 | 38.579 | nan | 50.369 | 30.941 |
[05/19 16:37:52 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/19 16:37:52 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP | category | AP | category | AP |
|:-----|:-----|:-----|:-----|:-----|:-----|
| headlamp | 35.879 | rear_bumper | 6.763 | door | 47.870 |

```

```
| hood          | 40.495 | front_bumper | 42.144 |          |          |
[05/19 16:37:52 d2.engine.defaults]: Evaluation results for car_part_val in csv format:
[05/19 16:37:52 d2.evaluation.testing]: cypypaste: Task: bbox
[05/19 16:37:52 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:37:52 d2.evaluation.testing]: cypypaste: 30.9477,52.8799,26.2704,nan,41.6017,27.9022
[05/19 16:37:52 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:37:52 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:37:52 d2.evaluation.testing]: cypypaste: 34.6302,50.8032,38.5791,nan,50.3686,30.9412
[05/19 16:37:52 d2.utils.events]: eta: 0:04:48 iter: 1999 total_loss: 0.6315 loss_cls: 0.1668 loss_box_reg: 0.2878 loss_mask: 0.1624 loss_rpn_cls: 0.004165 loss_rpn_loc: 0.0109 time: 0.5786 data_time: 0.0044 lr: 5e-05 max_mem: 1973M
[05/19 16:38:03 d2.utils.events]: eta: 0:04:36 iter: 2019 total_loss: 0.5999 loss_cls: 0.1694 loss_box_reg: 0.2643 loss_mask: 0.1311 loss_rpn_cls: 0.002982 loss_rpn_loc: 0.0114 time: 0.5785 data_time: 0.0048 lr: 2.5e-06 max_mem: 1973M
[05/19 16:38:15 d2.utils.events]: eta: 0:04:25 iter: 2039 total_loss: 0.6916 loss_cls: 0.1653 loss_box_reg: 0.2958 loss_mask: 0.1475 loss_rpn_cls: 0.003466 loss_rpn_loc: 0.01282 time: 0.5785 data_time: 0.0045 lr: 2.5e-06 max_mem: 1973M
[05/19 16:38:27 d2.utils.events]: eta: 0:04:13 iter: 2059 total_loss: 0.5947 loss_cls: 0.1579 loss_box_reg: 0.2775 loss_mask: 0.1609 loss_rpn_cls: 0.001282 loss_rpn_loc: 0.01237 time: 0.5787 data_time: 0.0041 lr: 2.5e-06 max_mem: 1973M
[05/19 16:38:38 d2.utils.events]: eta: 0:04:02 iter: 2079 total_loss: 0.555 loss_cls: 0.1594 loss_box_reg: 0.2472 loss_mask: 0.1333 loss_rpn_cls: 0.001669 loss_rpn_loc: 0.01134 time: 0.5788 data_time: 0.0040 lr: 2.5e-06 max_mem: 1973M
[05/19 16:38:50 d2.utils.events]: eta: 0:03:50 iter: 2099 total_loss: 0.6481 loss_cls: 0.1681 loss_box_reg: 0.2757 loss_mask: 0.1586 loss_rpn_cls: 0.002496 loss_rpn_loc: 0.01087 time: 0.5787 data_time: 0.0040 lr: 2.5e-06 max_mem: 1973M
[05/19 16:39:01 d2.utils.events]: eta: 0:03:38 iter: 2119 total_loss: 0.5258 loss_cls: 0.1478 loss_box_reg: 0.2369 loss_mask: 0.133 loss_rpn_cls: 0.002972 loss_rpn_loc: 0.01127 time: 0.5786 data_time: 0.0040 lr: 2.5e-06 max_mem: 1973M
[05/19 16:39:13 d2.utils.events]: eta: 0:03:27 iter: 2139 total_loss: 0.6017 loss_cls: 0.1519 loss_box_reg: 0.2813 loss_mask: 0.1432 loss_rpn_cls: 0.0008025 loss_rpn_loc: 0.01235 time: 0.5785 data_time: 0.0040 lr: 2.5e-06 max_mem: 1973M
[05/19 16:39:24 d2.utils.events]: eta: 0:03:15 iter: 2159 total_loss: 0.5686 loss_cls: 0.1733 loss_box_reg: 0.2636 loss_mask: 0.1305 loss_rpn_cls: 0.002543 loss_rpn_loc: 0.01119 time: 0.5784 data_time: 0.0038 lr: 2.5e-06 max_mem: 1973M
[05/19 16:39:36 d2.utils.events]: eta: 0:03:04 iter: 2179 total_loss: 0.6461 loss_cls: 0.1594 loss_box_reg: 0.2894 loss_mask: 0.1761 loss_rpn_cls: 0.005262 loss_rpn_loc: 0.01564 time: 0.5785 data_time: 0.0040 lr: 2.5e-06 max_mem: 1973M
[05/19 16:39:47 d2.utils.events]: eta: 0:02:52 iter: 2199 total_loss: 0.555 loss_cls: 0.1336 loss_box_reg: 0.2472 loss_mask: 0.1453 loss_rpn_cls: 0
```

.0006854 loss_rpn_loc: 0.01164 time: 0.5785 data_time: 0.0038 lr: 2.5e-06
max_mem: 1973M
[05/19 16:39:59 d2.utils.events]: eta: 0:02:41 iter: 2219 total_loss: 0.57
66 loss_cls: 0.168 loss_box_reg: 0.252 loss_mask: 0.1348 loss_rpn_cls: 0.
001924 loss_rpn_loc: 0.01417 time: 0.5784 data_time: 0.0037 lr: 2.5e-06
max_mem: 1973M
[05/19 16:40:10 d2.utils.events]: eta: 0:02:29 iter: 2239 total_loss: 0.57
89 loss_cls: 0.1534 loss_box_reg: 0.2831 loss_mask: 0.1505 loss_rpn_cls:
0.007307 loss_rpn_loc: 0.0147 time: 0.5783 data_time: 0.0039 lr: 2.5e-06
max_mem: 1973M
[05/19 16:40:21 d2.utils.events]: eta: 0:02:17 iter: 2259 total_loss: 0.69
54 loss_cls: 0.1608 loss_box_reg: 0.3095 loss_mask: 0.1469 loss_rpn_cls:
0.005565 loss_rpn_loc: 0.01313 time: 0.5782 data_time: 0.0037 lr: 2.5e-06
max_mem: 1973M
[05/19 16:40:33 d2.utils.events]: eta: 0:02:06 iter: 2279 total_loss: 0.57
9 loss_cls: 0.1492 loss_box_reg: 0.2755 loss_mask: 0.1301 loss_rpn_cls: 0
.001439 loss_rpn_loc: 0.01454 time: 0.5784 data_time: 0.0037 lr: 2.5e-06
max_mem: 1973M
[05/19 16:40:45 d2.utils.events]: eta: 0:01:54 iter: 2299 total_loss: 0.55
91 loss_cls: 0.1635 loss_box_reg: 0.2535 loss_mask: 0.1567 loss_rpn_cls:
0.0009827 loss_rpn_loc: 0.01301 time: 0.5786 data_time: 0.0039 lr: 2.5e-0
6 max_mem: 1973M
[05/19 16:40:57 d2.utils.events]: eta: 0:01:43 iter: 2319 total_loss: 0.53
59 loss_cls: 0.1546 loss_box_reg: 0.2433 loss_mask: 0.1364 loss_rpn_cls:
0.002618 loss_rpn_loc: 0.01046 time: 0.5786 data_time: 0.0038 lr: 2.5e-06
max_mem: 1973M
[05/19 16:41:09 d2.utils.events]: eta: 0:01:31 iter: 2339 total_loss: 0.57
59 loss_cls: 0.1673 loss_box_reg: 0.2731 loss_mask: 0.1443 loss_rpn_cls:
0.001201 loss_rpn_loc: 0.01203 time: 0.5787 data_time: 0.0040 lr: 2.5e-06
max_mem: 1973M
[05/19 16:41:21 d2.utils.events]: eta: 0:01:20 iter: 2359 total_loss: 0.61
78 loss_cls: 0.1466 loss_box_reg: 0.2872 loss_mask: 0.1597 loss_rpn_cls:
0.004857 loss_rpn_loc: 0.01733 time: 0.5788 data_time: 0.0040 lr: 2.5e-06
max_mem: 1973M
[05/19 16:41:32 d2.utils.events]: eta: 0:01:09 iter: 2379 total_loss: 0.57
3 loss_cls: 0.1553 loss_box_reg: 0.2319 loss_mask: 0.1307 loss_rpn_cls: 0
.001485 loss_rpn_loc: 0.0151 time: 0.5788 data_time: 0.0038 lr: 2.5e-06
max_mem: 1973M
[05/19 16:41:45 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:41:45 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:41:45 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...
[05/19 16:41:45 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:41:45 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:41:50 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0018 s/iter. Inference: 0.3082 s/iter. Eval: 0.1520 s/iter. Total: 0.4621 s
/iter. ETA=0:00:00
[05/19 16:41:50 d2.evaluation.evaluator]: Total inference time: 0:00:02.87082
1 (0.478470 s / iter per device, on 1 devices)

```

[05/19 16:41:50 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:01 (0.308245 s / iter per device, on 1 devices)
[05/19 16:41:50 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:41:50 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Evaluating predictions with u
nofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:41:51 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:41:51 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finishe
d in 0.04 seconds.
[05/19 16:41:51 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:41:51 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finis
hed in 0.04 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.3
40
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.5
31
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.3
29
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
13
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.3
23
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.4
06
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.5
34
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.5
40
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.
000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.4
55
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.5
72
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----|:-----|:-----|:-----|:-----|:-----|
| 34.003 | 53.125 | 32.918 | nan | 41.337 | 32.271 |
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Some metrics cannot be comput
ed and is shown as NaN.
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----|:-----|:-----|:-----|:-----|:-----|
| headlamp | 35.239 | rear_bumper | 9.193 | door | 43.313 |

```

```

| hood          | 43.911 | front_bumper | 38.359 |          |          |
Loading and preparing results...
DONE (t=0.02s)
creating index...
index created!
[05/19 16:41:51 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[05/19 16:41:51 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.04 seconds.
[05/19 16:41:51 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:41:51 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.04 seconds.
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.345
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.513
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.392
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.504
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.313
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.400
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.490
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.492
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.540
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.497
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Evaluation results for segm:
|  AP   | AP50 | AP75 | APs  | APm   | APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 34.508 | 51.254 | 39.222 | nan  | 50.392 | 31.283 |
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/19 16:41:51 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP   | category | AP   | category | AP   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 35.847 | rear_bumper | 6.763 | door      | 46.759 |
| hood     | 40.495 | front_bumper | 42.678 |           |         |
[05/19 16:41:51 d2.engine.defaults]: Evaluation results for car_part_val in csv format:
[05/19 16:41:51 d2.evaluation.testing]: copypaste: Task: bbox
[05/19 16:41:51 d2.evaluation.testing]: copypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:41:51 d2.evaluation.testing]: copypaste: 34.0030,53.1246,32.9182,nan,41.3368,32.2706

```

```
[05/19 16:41:51 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:41:51 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:41:51 d2.evaluation.testing]: cypypaste: 34.5084,51.2536,39.2221,na
n,50.3916,31.2825
[05/19 16:41:51 d2.utils.events]: eta: 0:00:57 iter: 2399 total_loss: 0.61
83 loss_cls: 0.1444 loss_box_reg: 0.2721 loss_mask: 0.1484 loss_rpn_cls: 0
.001132 loss_rpn_loc: 0.01297 time: 0.5788 data_time: 0.0039 lr: 2.5e-06
max_mem: 1973M
[05/19 16:42:03 d2.utils.events]: eta: 0:00:46 iter: 2419 total_loss: 0.66
97 loss_cls: 0.1661 loss_box_reg: 0.311 loss_mask: 0.1546 loss_rpn_cls: 0
.008387 loss_rpn_loc: 0.01424 time: 0.5791 data_time: 0.0044 lr: 2.5e-06
max_mem: 1973M
[05/19 16:42:15 d2.utils.events]: eta: 0:00:34 iter: 2439 total_loss: 0.59
41 loss_cls: 0.1529 loss_box_reg: 0.2742 loss_mask: 0.1447 loss_rpn_cls: 0
.001156 loss_rpn_loc: 0.01217 time: 0.5791 data_time: 0.0038 lr: 2.5e-06
max_mem: 1973M
[05/19 16:42:26 d2.utils.events]: eta: 0:00:23 iter: 2459 total_loss: 0.52
04 loss_cls: 0.1442 loss_box_reg: 0.2447 loss_mask: 0.1214 loss_rpn_cls: 0
.0012 loss_rpn_loc: 0.01345 time: 0.5790 data_time: 0.0038 lr: 2.5e-06
max_mem: 1973M
[05/19 16:42:38 d2.utils.events]: eta: 0:00:11 iter: 2479 total_loss: 0.65
82 loss_cls: 0.1742 loss_box_reg: 0.3016 loss_mask: 0.1681 loss_rpn_cls: 0
.002508 loss_rpn_loc: 0.01202 time: 0.5790 data_time: 0.0039 lr: 2.5e-06
max_mem: 1973M
[05/19 16:42:50 d2.utils.events]: eta: 0:00:00 iter: 2499 total_loss: 0.59
7 loss_cls: 0.1434 loss_box_reg: 0.265 loss_mask: 0.1635 loss_rpn_cls: 0.
002852 loss_rpn_loc: 0.009996 time: 0.5789 data_time: 0.0041 lr: 2.5e-06
max_mem: 1973M
[05/19 16:42:50 d2.engine.hooks]: Overall training speed: 2498 iterations in
0:24:06 (0.5789 s / it)
[05/19 16:42:50 d2.engine.hooks]: Total training time: 0:25:04 (0:00:58 on ho
oks)
[05/19 16:42:50 d2.data.datasets.coco]: Loaded 11 images in COCO format from
damage_dataset/val/COCO_mul_val_annos.json
[05/19 16:42:50 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used i
n inference: [ResizeShortestEdge(short_edge_length=(800, 800), max_size=1333,
sample_style='choice')]
[05/19 16:42:50 d2.data.common]: Serializing 11 elements to byte tensors and
concatenating them all ...
[05/19 16:42:50 d2.data.common]: Serialized dataset takes 0.01 MiB
[05/19 16:42:50 d2.evaluation.evaluator]: Start inference on 11 batches
[05/19 16:42:56 d2.evaluation.evaluator]: Inference done 11/11. Dataloading:
0.0018 s/iter. Inference: 0.3048 s/iter. Eval: 0.1592 s/iter. Total: 0.4657 s
/iter. ETA=0:00:00
[05/19 16:42:56 d2.evaluation.evaluator]: Total inference time: 0:00:02.91283
1 (0.485472 s / iter per device, on 1 devices)
[05/19 16:42:56 d2.evaluation.evaluator]: Total inference pure compute time:
0:00:01 (0.304769 s / iter per device, on 1 devices)
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Preparing results for COCO fo
rmat ...
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Saving results to coco_eval/c
oco_instances_results.json
```

```

[05/19 16:42:56 d2.evaluation.coco_evaluation]: Evaluating predictions with unofficial COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
[05/19 16:42:56 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[05/19 16:42:56 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.04 seconds.
[05/19 16:42:56 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:42:56 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.05 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.328
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.526
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.404
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.311
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.528
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.534
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.445
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.569
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 32.779 | 52.591 | 31.844 | nan | 40.371 | 31.089 |
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 35.169 | rear_bumper | 7.176 | door | 39.921 |
| hood | 44.522 | front_bumper | 37.109 | | |
Loading and preparing results...
DONE (t=0.02s)
creating index...
index created!
[05/19 16:42:56 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*

```



```

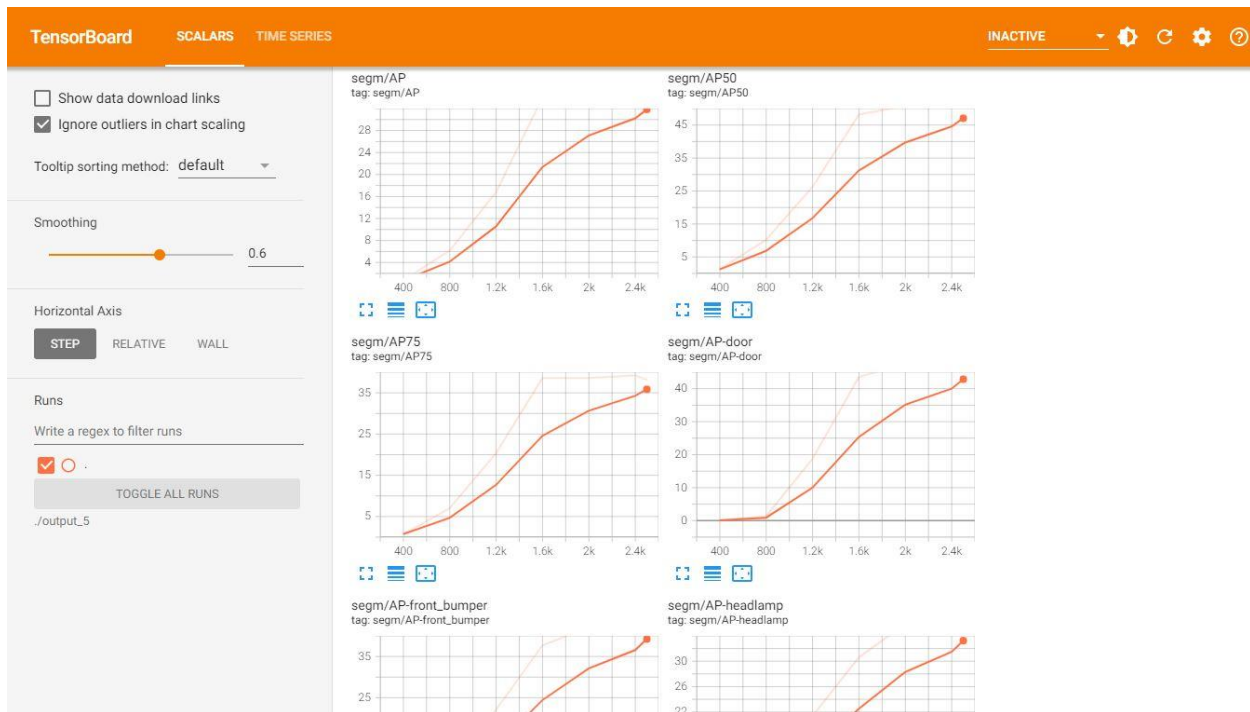
[05/19 16:42:56 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.04 seconds.
[05/19 16:42:56 d2.evaluation.fast_eval_api]: Accumulating evaluation results
...
[05/19 16:42:56 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.03 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.342
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.507
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.382
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.504
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.304
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.403
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.490
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.492
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.540
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.497
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 34.190 | 50.657 | 38.152 | nan | 50.387 | 30.425 |
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[05/19 16:42:56 d2.evaluation.coco_evaluation]: Per-category segm AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| headlamp | 35.793 | rear_bumper | 4.609 | door | 46.983 |
| hood | 40.495 | front_bumper | 43.070 | | |
[05/19 16:42:56 d2.engine.defaults]: Evaluation results for car_part_val in csv format:
[05/19 16:42:56 d2.evaluation.testing]: cypypaste: Task: bbox
[05/19 16:42:56 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:42:56 d2.evaluation.testing]: cypypaste: 32.7794,52.5907,31.8439,nan,40.3710,31.0894
[05/19 16:42:56 d2.evaluation.testing]: cypypaste: Task: segm
[05/19 16:42:56 d2.evaluation.testing]: cypypaste: AP,AP50,AP75,APs,APm,APl
[05/19 16:42:56 d2.evaluation.testing]: cypypaste: 34.1899,50.6567,38.1518,nan,50.3871,30.4249

```

In []:


```
%reload_ext tensorboard
%tensorboard --logdir ./output_2
```

Average Precision vs Epochs



Loss vs Epochs





Predictions of car part detection model

Preparing test data

As seen from above model **2** detects two moderate damages and model **3** detects 4 car parts among which **front bumper** segmentation is **overlapping** with one **damage** segmentation. Hence we calculate **dice coefficient** and put it under “front_bumper_dice” column.

Another damage in the front fender (beside headlamp) does not overlap any segmentation from the second model because the training data did not have any annotations for the **fender** of any car. Thus it will go to the unknown column.

Test data of the above image

Final prediction of price

Now as mentioned earlier the training dataset was **too small** to fit a model so instead of fitting a model (supposedly model **1**), for every row in test data I calculated mean price of identical rows in train data. In the end I added all the estimated prices from each row to get the total price.

Future work

- A large dataset would help us train an actual machine learning model for model **1** instead of the calculating mean as we did above.
- Proper annotation of all external parts of car will improve model **3** performance and we won't need the 'unknown' column in the train or test data. (Eg. In the above test example if the front fender of cars were annotated in train data, overlap would've happened between the damage and the fender segmentation outputs from the two models **2** & **3** respectively and we would've filled the 'fender_dice' column instead of the unknown column).

- Proper price estimates will make the model **1** outputs reliable and such data is only present with the insurance company.
- In reality price of repairing a car does not only depend on the damages but on many other factors too like - **brand of car, age, paint color, model number** & so on. These too are available to the insurance company and we can add these as features to the train data to improve model performance.
- We can also use **IOU** score to find the percentage of overlap between two segmentation masks.
- Segmented damaged area images can also be added to the training dataset for every row which would make use of **CNN+MLP** model to train model **1**.
- It only makes sense to find the percentage overlap using dice coefficient if the damage as well as the whole car part where the damage has occurred is visible (Eg. In the above test example the whole front bumper of the car is not visible, therefore the dice coefficient between the damage and the front bumper doesn't make much sense).

Deployment

The code can be found [here](#). I have deployed the model using Flask framework. Follow the steps on GitHub to deploy the model in your local machine. This is a short video of my web-app:

Deployed w